

Small Image Classification Using Convolutional Neural Network (CNN)

In this notebook, we will classify small images cifar10 dataset from tensorflow keras datasets. There are total 10 classes as shown below. We will use CNN for classification



In [1]:

```
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
import numpy as np
Load the dataset
```

```
In [2]:
(X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()
X_train.shape
Out[2]:
(50000, 32, 32, 3)
In [3]:
X_test.shape
Out[3]:
(10000, 32, 32, 3)
Here we see there are 50000 training images and 1000 test images
```

```
In [4]:
y_train.shape
Out[4]:
(50000, 1)
In [5]:
y_train[:5]
Out[5]:
array([[6],
       [9],
       [9],
       [4],
       [1]], dtype=uint8)
y_train is a 2D array, for our classification having 1D array is good enough. so we will convert
this to now 1D array
```

```
In [6]:
y_train = y_train.reshape(-1,)
y_train[:5]
Out[6]:
array([6, 9, 9, 4, 1], dtype=uint8)
In [7]:
y_test = y_test.reshape(-1,)
In [8]:
classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
Let's plot some images to see what they are
```

```
In [9]:
def plot_sample(X, y, index):
    plt.figure(figsize = (15,2))
    plt.imshow(X[index])
    plt.xlabel(classes[y[index]])
In [10]:
plot_sample(X_train, y_train, 0)
```

In [11]:

```
plot_sample(X_train, y_train, 1)
```

Normalize the images to a number from 0 to 1. Image has 3 channels (R,G,B) and each value in the channel can range from 0 to 255. Hence to normalize in 0-->1 range, we need to divide it by 255

Normalizing the training data

In [12]:

```
X_train = X_train / 255.0
```

```
X_test = X_test / 255.0
```

Build simple artificial neural network for image classification

In [13]:

```
ann = models.Sequential([
    layers.Flatten(input_shape=(32,32,3)),
    layers.Dense(3000, activation='relu'),
    layers.Dense(1000, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

```
ann.compile(optimizer='SGD',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])
```

```
ann.fit(X_train, y_train, epochs=5)
```

Epoch 1/5

```
1563/1563 [=====] - 2s 2ms/step - loss: 1.8074 - accuracy: 0.3561
```

Epoch 2/5

```
1563/1563 [=====] - 2s 1ms/step - loss: 1.6208 - accuracy: 0.4285
```

Epoch 3/5

```
1563/1563 [=====] - 2s 2ms/step - loss: 1.5380 - accuracy: 0.4585
```

Epoch 4/5

```
1563/1563 [=====] - 2s 2ms/step - loss: 1.4808 - accuracy: 0.4806
```

Epoch 5/5

```
1563/1563 [=====] - 2s 2ms/step - loss: 1.4326 - accuracy: 0.4928
```

Out[13]:

```
<tensorflow.python.keras.callbacks.History at 0x295ab873c10>
```

You can see that at the end of 5 epochs, accuracy is at around 49%

In [14]:

```
from sklearn.metrics import confusion_matrix , classification_report
```

```
import numpy as np
```

```
y_pred = ann.predict(X_test)
```

```
y_pred_classes = [np.argmax(element) for element in y_pred]
```

```
print("Classification Report: \n", classification_report(y_test, y_pred_classes))
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.63 | 0.45 | 0.53 | 1000 |
| 1 | 0.72 | 0.46 | 0.56 | 1000 |
| 2 | 0.33 | 0.46 | 0.39 | 1000 |
| 3 | 0.36 | 0.25 | 0.29 | 1000 |
| 4 | 0.44 | 0.37 | 0.40 | 1000 |
| 5 | 0.34 | 0.46 | 0.39 | 1000 |
| 6 | 0.56 | 0.47 | 0.51 | 1000 |
| 7 | 0.39 | 0.67 | 0.50 | 1000 |
| 8 | 0.64 | 0.60 | 0.62 | 1000 |
| 9 | 0.59 | 0.53 | 0.55 | 1000 |
| accuracy | | | 0.47 | 10000 |
| macro avg | 0.50 | 0.47 | 0.47 | 10000 |
| weighted avg | 0.50 | 0.47 | 0.47 | 10000 |

Now let us build a convolutional neural network to train our images

In [15]:

```
cnn = models.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```

In [16]:

```
cnn.compile(optimizer='adam',
            loss='sparse_categorical_crossentropy',
            metrics=['accuracy'])
```

In [17]:

```
cnn.fit(X_train, y_train, epochs=10)
Epoch 1/10
1563/1563 [=====] - 2s 2ms/step - loss: 1.4407 - accuracy:
0.4810
Epoch 2/10
1563/1563 [=====] - 2s 2ms/step - loss: 1.1084 - accuracy:
0.6109
Epoch 3/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.9895 - accuracy:
0.6574
Epoch 4/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.9071 - accuracy:
0.6870
Epoch 5/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.8416 - accuracy:
0.7097
Epoch 6/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.7847 - accuracy:
0.7262
Epoch 7/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.7350 - accuracy:
0.7448
Epoch 8/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.6941 - accuracy:
0.7574
Epoch 9/10
1563/1563 [=====] - 2s 1ms/step - loss: 0.6516 - accuracy:
0.7731
Epoch 10/10
1563/1563 [=====] - 2s 2ms/step - loss: 0.6187 - accuracy:
0.7836
```

Out[17]:

```
<tensorflow.python.keras.callbacks.History at 0x296555783d0>
```

With CNN, at the end 5 epochs, accuracy was at around 70% which is a significant improvement over ANN. CNN's are best for image classification and gives superb accuracy. Also computation is much less compared to simple ANN as maxpooling reduces the image dimensions while still preserving the features

In [18]:

```
cnn.evaluate(X_test,y_test)
313/313 [=====] - 0s 1ms/step - loss: 0.9022 - accuracy:
0.7028
```

Out[18]:

```
[0.9021560549736023, 0.7027999758720398]
```

```
In [19]:
y_pred = cnn.predict(X_test)
y_pred[:5]
Out[19]:
array([[4.3996371e-04, 3.4844263e-05, 1.5558505e-03, 8.8400185e-01,
        1.9452239e-04, 3.5314459e-02, 7.2777577e-02, 6.9044131e-06,
        5.6417785e-03, 3.2224660e-05],
       [8.1062522e-03, 5.0841425e-02, 1.2453231e-07, 5.3348430e-07,
        9.1728407e-07, 1.0009186e-08, 2.8985988e-07, 1.7532484e-09,
        9.4089705e-01, 1.5346886e-04],
       [1.7055811e-02, 1.1841061e-01, 4.6799007e-05, 2.7727904e-02,
        1.0848254e-03, 1.0896578e-03, 1.3575243e-04, 2.8652203e-04,
        7.8895986e-01, 4.5202184e-02],
       [3.1300801e-01, 1.1591638e-02, 1.1511055e-02, 3.9592334e-03,
        7.7280165e-03, 5.6289224e-05, 2.3531138e-04, 9.4204297e-06,
        6.5178138e-01, 1.1968113e-04],
       [1.3230885e-05, 2.1221960e-05, 9.2594400e-02, 3.3585075e-02,
        4.4722903e-01, 4.1028224e-03, 4.2241842e-01, 2.8064171e-05,
        6.6392668e-06, 1.0745022e-06]], dtype=float32)
```

```
In [20]:
y_classes = [np.argmax(element) for element in y_pred]
y_classes[:5]
Out[20]:
[3, 8, 8, 8, 4]
```

```
In [21]:
y_test[:5]
Out[21]:
array([3, 8, 8, 0, 6], dtype=uint8)
In [22]:
plot_sample(X_test, y_test, 3)
```

```
In [23]:
classes[y_classes[3]]
```

```
Out[23]:
'ship'
```

```
In [24]:
classes[y_classes[3]]
```

```
Out[24]:
'ship'
```